

# IRWA Project Part 4: RAG, User Interface, and Web Analytics

---

Authors: Lucas Andreu, Pau Chaves, Pol Bonet, Joan Company

Course: Information Retrieval and Web Analytics (IRWA)

Date: November 29, 2025

## 1. Introduction

In Part 4 we turn the search engine from Parts 1–3 into a full web application with:

- A Flask-based UI for querying the product corpus and exploring individual products.
- A Retrieval-Augmented Generation (RAG) component that summarizes and recommends products using an external LLM (Groq API).
- A web-analytics layer that tracks how users interact with the system (requests, queries, clicks, dwell time, sessions, and user context) and visualizes this data in a dashboard.

All analytics are stored in memory (Python objects) for easy reproduction by the instructors.

## 2. User Interface & Search Integration

### 2.1 Search Page

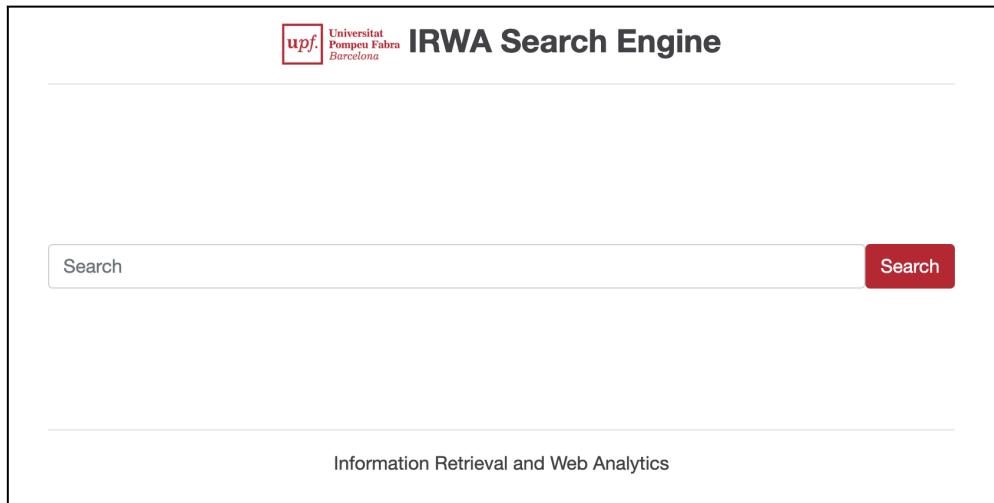
The main search page is implemented in `templates/index.html` and served by the `/` route in `web_app.py`. It contains:

- A central search box (`<input name="search-query">`).
- A Search button submitting the form via `POST` to `/search`.

HTML

```
<form class="d-flex" method="POST" onSubmit='return validate();'  
action="/search">  
    <input class="form-control me-2" name="search-query" type="search"  
placeholder="Search" autofocus="autofocus">  
    <button class="btn btn-primary" type="submit">Search</button>  
</form>
```

The layout uses Bootstrap and a shared `base.html` template for header/footer.



**Figure 1** – Home / search page, before submitting any query

## 2.2 Search Action and Engine Integration

The `/search` route in `web_app.py`:

```
Python
@app.route('/search', methods=['POST'])
def search_form_post():
    search_query = request.form['search-query']
    ...
    search_id = analytics_data.save_query_terms(
        terms=search_query,
        ip=request.remote_addr,
        user_agent=request.headers.get("User-Agent", ""),
        browser=request.user_agent.browser,
        session_id=session.get("session_id")
    )
    ...
    results = search_engine.search(search_query, search_id, corpus)
```

Key points:

- The raw query string is passed to `SearchEngine.search(search_query, search_id, corpus)`.
- A `search_id` is generated and stored; this links queries to later clicks (for analytics).
- Query metadata (IP, browser, session, terms, etc.) is registered via `AnalyticsData.save_query_terms`.

The **general search function** lives in `myapp/search/search_engine.py`:

Python

```
class SearchEngine:  
    def search(self, search_query, search_id, corpus):  
        results = search_in_corpus(  
            query=search_query,  
            search_id=search_id,  
            corpus=corpus,  
            method="bm25",  
            k=20,  
            use_and=True  
        )  
        return results
```

This function only needs a string query (plus `search_id` for logging) and can be easily reused from any UI.

## 2.3 Search Algorithms

The core retrieval logic is in `myapp/search/algorithms.py`:

- **Preprocessing and indexing**
  - We load an enriched Flipkart fashion dataset and a boolean inverted index (`fashion_products_dataset_enriched.json`, `boolean_inverted_index.json`).
  - Index fields: `title_clean`, `description_clean`, `metadata_clean`.
- **Candidate generation**
  - AND semantics by default: `_candidate_docs_and(q_terms)` intersects postings lists for all query terms.
  - Fallback to OR semantics if AND yields no candidates.
- **Ranking algorithms implemented**
  - **BM25** (default):
    - Standard Okapi BM25 with document length normalization and pre-computed `idf_bm25`.
  - **TF-IDF cosine similarity**:
    - `tfidf_weights` and `doc_norms` pre-computed for fast cosine scoring.
  - **Custom TF-IDF + business boost**:
    - `_numeric_boost` favors in-stock products with better ratings, higher discounts and reasonable prices.
    - `scores = tfidf_score * _numeric_boost(record)`.
- **Result representation**
  - `search_in_corpus(...)` converts each hit into a `ResultItem` object (Pydantic model) with:

- `title, description`
- `selling_price, actual_price, discount, average_rating, out_of_stock`
- `url` → internal `/doc_details?pid=...&search_id=...`
- `source_url` → original Flipkart URL
- `ranking` → BM25 (or other) score, used later in analytics.

This structure is suitable for a web application: the search engine is pure Python (no Flask dependency) and returns serializable objects.

## 2.4 Results Page

The `/search` route renders `templates/results.html` with:

- The **number of results** found.
- An **AI-Generated Summary** (RAG, described in Section 3).
- A list of ranked results using the metadata specified in the assignment:

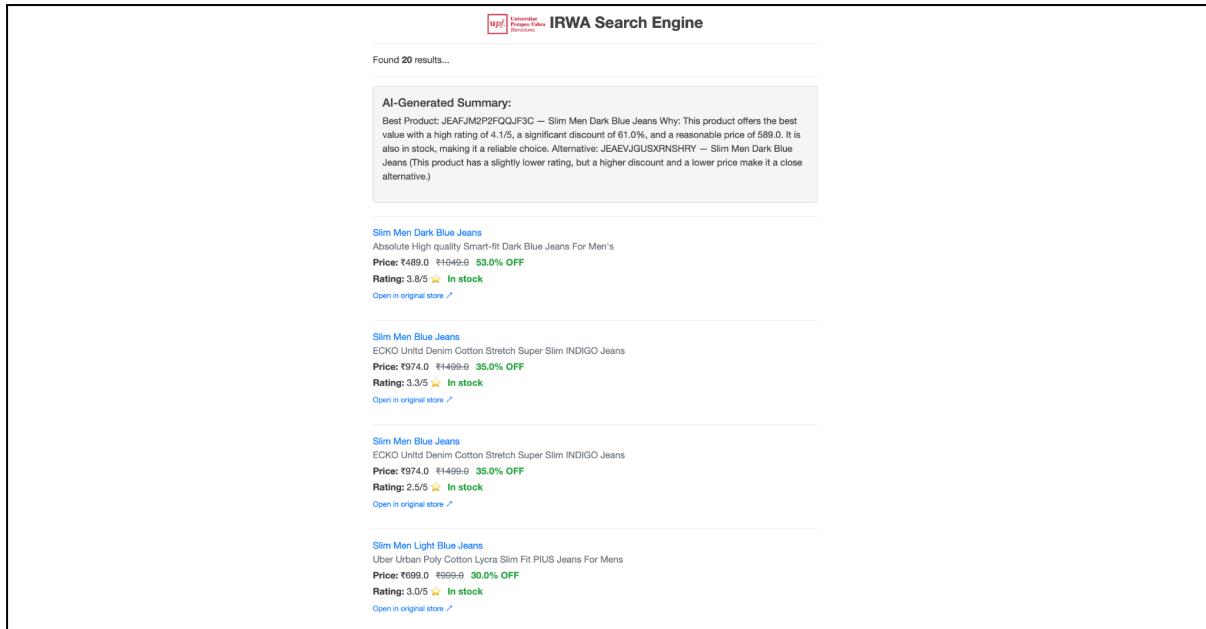
HTML

```
<div class="doc-title">
    <a href="{{ item.url }}">{{ item.title }}</a>
</div>
<div class="doc-desc text-muted">{{ item.description }}</div>

<div class="mt-1">
    <span><strong>Price:</strong> ₹{{ item.selling_price }}</span>
    <span class="ml-2 text-muted"><del>₹{{ item.actual_price }}</del></span>
    <span class="ml-2 text-success"><strong>{{ item.discount }}% OFF</strong></span>
</div>

<div class="mt-1">
    <span><strong>Rating:</strong> {{ item.average_rating }}/5 ★</span>
    {% if item.out_of_stock %}
        <span class="ml-2 text-danger"><strong>Out of stock</strong></span>
    {% else %}
        <span class="ml-2 text-success"><strong>In stock</strong></span>
    {% endif %}
</div>

<a href="{{ item.source_url }}" target="_blank" class="small">
    Open in original store ↗
</a>
```



**Figure 2 – Results Page Showing Ranked Product List (men slim blue jeans)**

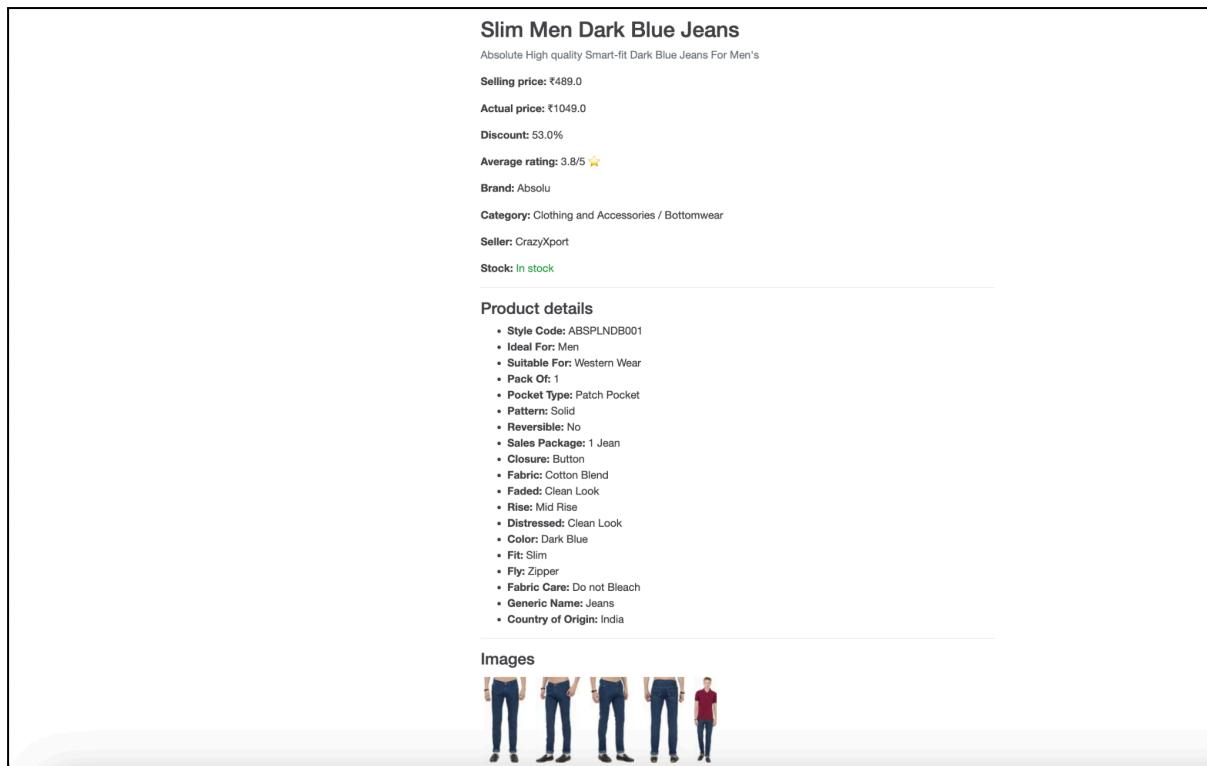
## 2.5 Document Details Page

The `/doc_details` route:

- Receives `pid` and `search_id`.
- Looks up the `Document` from the loaded `corpus`.
- Registers a click event (with ranking, query, and user context) in `AnalyticsData.register_click`.
- Stores `last_click_time` in the session so that the next search can compute dwell time.

`templates/doc_details.html` shows:

- Full title and description.
- Price, discount and rating.
- Brand, category & sub-category, seller.
- Stock status.
- A formatted list of `product_details` (technical specs).
- Product images if available.
- A link to the original Flipkart product.



**Figure 3 – All info details page (men slim blue jeans)**

### 3. RAG: Retrieval-Augmented Generation

#### 3.1 Baseline RAG and limitations

The instructor's baseline RAG pipeline:

- Passed only a list of product PIDs and titles to the Groq LLM.
- Used a generic prompt without emphasis on stock status or value.
- Let the model freely choose any product, which sometimes:
  - Ignored stock, discount, or rating.
  - Fabricated attributes.
  - Returned ambiguous recommendations.
  - Always chose a product, even if nothing matched the query.

#### 3.2 Implemented improvements

The improved RAG logic is implemented in `myapp/generation/rag.py` (`RAGGenerator`).

##### 1) Richer metadata + helper score

We pass a structured list of retrieved products to the LLM using `_format_results`:

Python

```
f" {i}. PID: {pid}\n"
f"   Title: {title}\n"
f"   Brand: {brand}\n"
f"   Category: {category} / {subcat}\n"
f"   Price: {price} | Actual: {actual_price} | Discount: {discount}%\n"
f"   Rating: {rating}/5 | In stock: {not bool(stock)}\n"
f"   retrieval_score: {retrieval_score} | helper_score: {helper}\n"
```

The `helper_score` is a deterministic value score:

Python

```
rating_norm    # normalized 0-1
discount_norm # normalized discount
price_norm     # cheaper is better
```

These fields guide the LLM toward products that are both **relevant and good value**, without hard-coding a ranking inside the model.

### Effect in practice

For a query like `men slim blue jeans`, the summary typically selects a **high-discount, good-rating, in-stock** pair of jeans and explicitly references price, discount and rating in the explanation.

## 2) Pre-filter obviously bad candidates

Before calling the LLM we filter out out-of-stock products if there are enough remaining hits:

Python

```
in_stock = [r for r in retrieved_results if not getattr(r, "out_of_stock",
False)]
if len(in_stock) >= 3:
    retrieved_results = in_stock
```

This prevents recommending unavailable products and simplifies the LLM's job by focusing on viable options.

### Effect in practice

For queries with both in-stock and out-of-stock items (e.g., `women solid cotton kurta beige`), the LLM now recommends an in-stock product with good discount and rating instead of a cheaper but unavailable alternative.

### 3) Safer prompting, low temperature, and “no good products” handling

We add a restrictive system prompt:

```
Python
SYSTEM_PROMPT = (
    "You are a careful, non-hallucinating product recommender. "
    "You must only use the provided retrieved products. "
    "If the retrieved products do not fit, say so explicitly."
)
```

The user prompt instructs the model to:

- Use the metadata to justify the recommendation.
- Prefer in-stock items.
- Output a fixed format:
  - Best Product: <PID> — <Title>
  - Why: ...
  - Alternative (optional): ...
- If nothing fits, output exactly:  
"There are no good products that fit the request based on the retrieved results."

We also reduce creativity with:

```
Python
temperature=0.1,
max_tokens=300
```

and add a simple fallback if the model returns an empty string.

#### Effect in practice

- For a realistic query like **winter boots waterproof women**, the answer is:  
"Best Product: SHOF97KQWHH7FDTG — Boots For Women (Beige)  
Why: This product is the best match for the user's request as it is a pair of boots for women, which aligns with the query intent. It is also waterproof..."
- For intentionally mismatched queries such as **neon green tuxedo** or **kids shoes**, the system now correctly returns:  
"There are no good products that fit the request based on the retrieved results."

### AI-Generated Summary:

Best Product: SHOF97KQWHH7FDTG — Boots For Women (Beige) Why: This product is the best match for the user's request as it is a pair of boots for women, which aligns with the query intent. Additionally, it is waterproof, making it a suitable choice for winter. The product is also in stock and has a good rating of 4.2/5. Alternative: None

**Figure 4** – RAG Summary for a Well-Matched Query (“winter boots waterproof women”)

### AI-Generated Summary:

There are no good products that fit the request based on the retrieved results.

**Figure 5** – RAG Output for a Query with No Suitable Products (“neon green tuxedo”)

## 4. Web Analytics

Analytics are implemented in `myapp/analytics/analytics_data.py` and integrated via `web_app.py`.

All data is stored in memory as Python lists/dicts so instructors can reproduce results simply by running the app and interacting with it.

### 4.1 Data collection

We collect four main event types.

#### 4.1.1 HTTP requests

`@app.before_request` in `web_app.py` logs every request:

```
Python
analytics_data.register_request(
    path=request.path,
    method=request.method,
    user_agent=request.headers.get("User-Agent", ""),
    ip=request.remote_addr,
    session_id=session.get("session_id"),
    ts=time.time()
)
```

Stored fields (table `requests`):

- `ts` – timestamp.

- `path` – URL path ("/", "/search", "/doc\_details", etc.).
- `method` – HTTP method (GET/POST).
- `user_agent` – full user agent string.
- `ip` – remote IP address.
- `session_id` – synthetic ID stored in the Flask session cookie (used to group user journeys).

#### 4.1.2 Queries

When `/search` is called, we generate a `search_id` and save the query via:

```
Python
search_id = analytics_data.save_query_terms(
    terms=search_query,
    ip=request.remote_addr,
    user_agent=request.headers.get("User-Agent", ""),
    browser=request.user_agent.browser,
    session_id=session.get("session_id")
)
```

Under the hood, `register_query` records:

- `ts, search_id, query`
- `n_terms` and tokenized `terms`
- User context:
  - `ip`
  - `user_agent`
  - `browser` (parsed from Flask's `request.user_agent.browser`)
  - `session_id`

This satisfies the requirement to collect *number of terms*, *order*, and *user context* (IP, browser, time of day via `ts`).

#### 4.1.3 Clicks and ranking

In `/doc_details` we link each click to the query that produced it:

```
Python
analytics_data.register_click(
    pid=pid,
    search_id=int(search_id) if search_id else -1,
    rank=getattr(doc_obj, "ranking", None),
    query=session.get("last_search_query"),
    ip=request.remote_addr,
    user_agent=request.headers.get("User-Agent", ""),
    session_id=session.get("session_id")
```

)

Stored fields (`clicks` table):

- `ts, pid, search_id`
- `rank` – BM25/TF-IDF score of the clicked document.
- `query` – last search query from the same session.
- `ip, user_agent, session_id`

We also maintain `fact_clicks[pid]` for quick “top clicked products” statistics.

The screenshot shows the 'IRWA Search Engine' interface with a 'Quick Stats' section. It lists several products with their IDs and brief descriptions:

- (2 visits) — id: JEAEVJGUSXRNSHRY — Slim Men Dark Blue Jeans - Absolute High quality Smart-fit Dark Blue Jeans For Men's
- (1 visits) — id: KTAZM3F8GGGPZC4 — Women Solid Cotton Blend Straight Kurta (Beige) -
- (1 visits) — id: RNCFXRFSHREUUZCN — Solid Men Raincoat -
- (1 visits) — id: RNCF4BYMQD53WHSF — Solid Men Raincoat - Type: Rainsuit a rain coat jackets & a pant for men women kids Material: Polyester sweater Sleeve: Full Sleeve mens rainwear Drawstring hooded neck For men Closure: Zip Colour: NeBlack
- (1 visits) — id: JEAFSKYHRVZSABPR — Slim Men Blue Jeans - ECKO Unltd Denim Cotton Stretch Super Slim INDIGO Jeans
- (1 visits) — id: KTAFX3FUTMB5HYXB — Women Solid Pure Cotton Straight Kurta (Beige) -

Information Retrieval and Web Analytics

**Figure 5 – Top Clicked Products Page (Stats)**

#### 4.1.4 Dwell time

We measure *dwell time* as the time between clicking a result and issuing the next search:

- On click (`/doc_details`), we save:

```
Python
session["last_click_time"] = time.time()
session["last_clicked_pid"] = pid
session["last_search_id"] = int(search_id) if search_id else -1
```

- On the next `/search`, before processing the new query:

```
Python
dwell = time.time() - session["last_click_time"]
analytics_data.register_dwell(
    pid=session["last_clicked_pid"],
    search_id=session.get("last_search_id", -1),
    dwell_seconds=dwell
)
```

Stored fields (`dwell_times` table):

- `ts, pid, search_id, dwell_seconds.`

This completes query→click→dwell linkage for later funnel analysis.

## 4.2 Data model (in-memory star schema)

Conceptually we follow a small **star schema**:

- **Fact tables**
  - `queries`: one row per issued query (with `search_id`).
  - `clicks`: one row per click on a result.
  - `dwell_times`: one row per interaction with dwell measurement.
  - `requests`: one row per HTTP request to any endpoint.
- **Dimensions / keys**
  - `search_id` links queries, clicks and dwell records.
  - `pid` links clicks/dwell to the **Document** (product) in the corpus.
  - `session_id` groups multiple queries and clicks into a physical session (user journey).
  - `ts` (timestamp) enables temporal analysis such as time-of-day.

This schema is implemented purely in Python lists/dicts but could be easily translated to relational tables if a database were used.

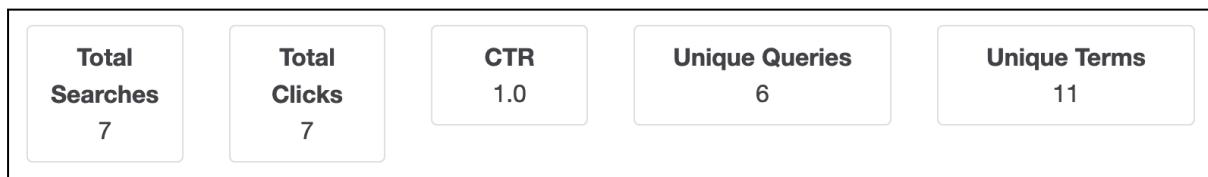
## 4.3 Key metrics and reports

### 4.3.1 Summary statistics

`AnalyticsData.summary_stats()` computes:

- `total_searches` – number of queries.
- `total_clicks` – number of clicks.
- `ctr` – click-through-rate = `clicks / searches`.
- `unique_queries` – distinct query strings.
- `unique_terms` – distinct terms across all queries.
- `avg_dwell` – average dwell time in seconds.

These are shown as KPI cards at the top of `dashboard.html`.



**Figure 6 – Dashboard High-Level KPIs**

#### 4.3.2 Funnel metrics (search → click → dwell)

`AnalyticsData.funnel_metrics()` further computes:

- `searches, clicks`
- `dwell_over_5s` – number of interactions where `dwell_seconds ≥ 5`.
- `ctr` – as above.
- `engagement_rate` – fraction of searches that led to a click with dwell  $\geq 5$ s.

These KPIs highlight how many searches turn into engaged product views (not just accidental clicks).

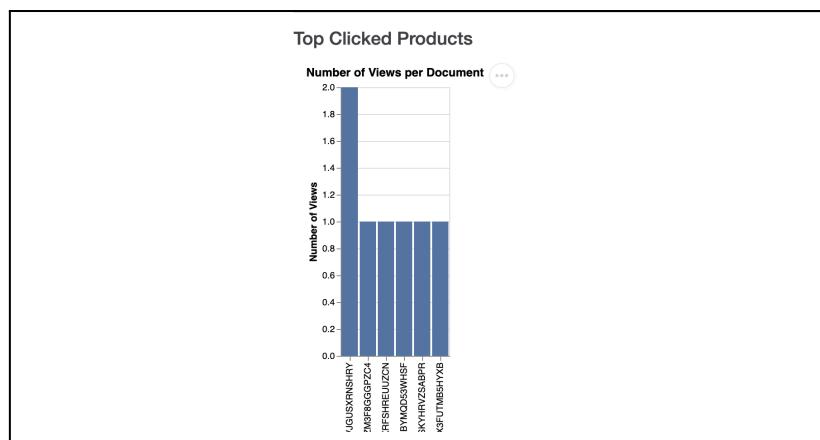


**Figure 7 – Dashboard KPIs (further)**

#### 4.3.3 Top clicked products

`plot_number_of_views()`:

- Uses `fact_clicks` to plot a bar chart “Number of Views per Document”.
- Embedded via an `<iframe>` /`plot_number_of_views` in the dashboard.

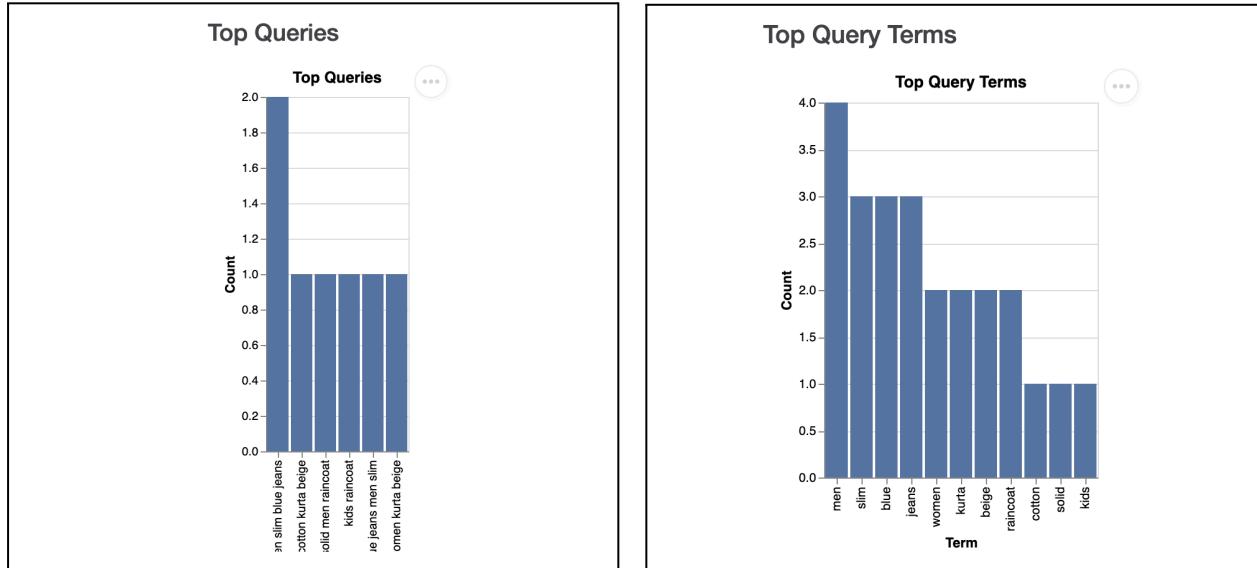


**Figure 8 – Bar Chart Number of Views**

#### 4.3.4 Query-level insights

- `plot_top_queries()` – bar chart of most frequent full queries.
- `plot_top_terms()` – bar chart of most frequent individual terms.

Both are embedded as iframes in `dashboard.html`.



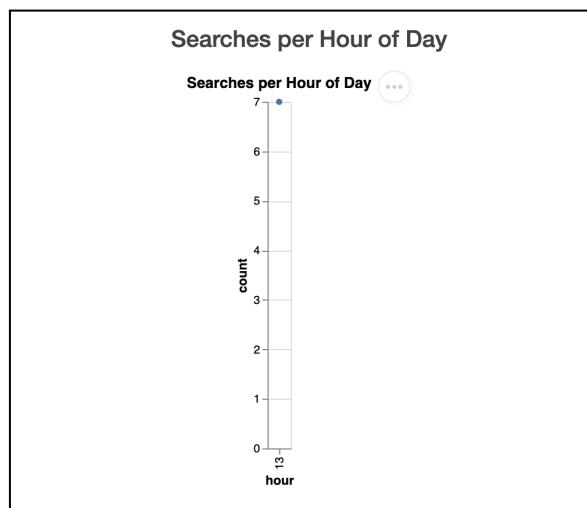
**Figure 9** – Bar Chart Top Queries and Top Query Terms

#### 4.3.5 Temporal analysis: searches per hour

`plot_searches_per_hour()`:

- Converts query timestamps `ts` to hour-of-day ("00"–"23").
- Groups by hour and plots a line chart “Searches per Hour of Day”.

This would show, for a real deployment, when the search engine is most used.



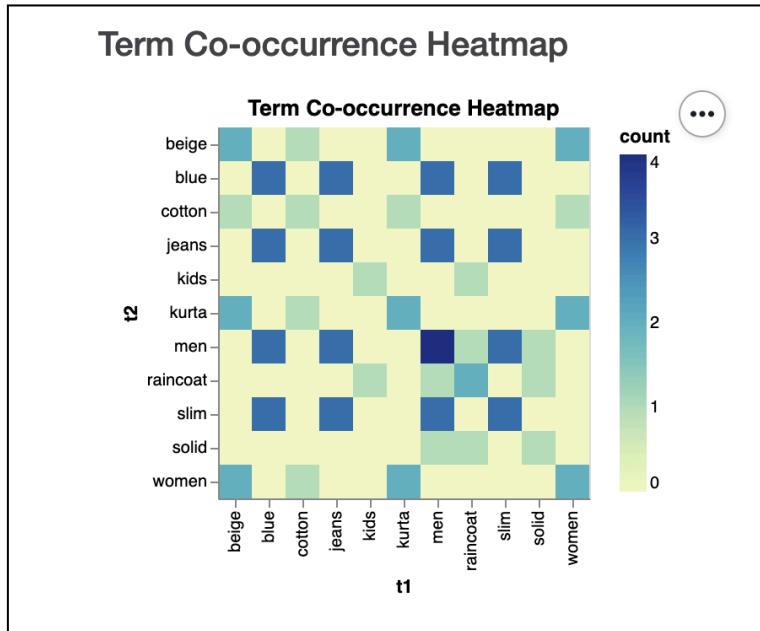
**Figure 10** – Searches per hour (all queries were searched at 1pm)

#### 4.3.6 Term co-occurrence heatmap (bonus)

`plot_term_heatmap(k=20):`

- Finds top-k most frequent terms across queries.
- Builds a co-occurrence matrix over unique terms in each query.
- Visualizes as an Altair heatmap (term vs term, colored by co-occurrence count).

This reveals which terms tend to be queried together (e.g., men + slim + jeans).



**Figure 11 – Term Co-occurrence Heatmap**

#### 4.3.7 Session paths (user journeys) – bonus

`session_paths()`:

- Groups requests by `session_id`.
- For each session outputs the sequence of paths visited (e.g., `/ → /search → /doc_details → /search`).

Displayed as a list under “User Journeys (Session Paths)” in `dashboard.html`. This approximates **time-based sessions**.

## **Figure 12 – User Journeys**

#### 4.3.8 Intent clusters (missions) – bonus

intent\_clusters():

- For each query, normalizes its set of terms into a sorted string key.
  - Counts how many times each unique term set appears.
  - Sorted by descending frequency.

Displayed as “Query Intent Clusters” with entries like:

3 searches — *blue jeans men slim*  
2 searches — *cotton kurta women beige*

This approximates **logical missions**: repeated attempts to satisfy the same information need with slightly different queries.

# Query Intent Clusters

- 3 searches — blue jeans men slim
  - 1 searches — beige cotton kurta women
  - 1 searches — men raincoat solid
  - 1 searches — kids raincoat
  - 1 searches — beige kurta women

**Figure 13** – Query Intent Clusters

## 4.4 Dashboard UI

The `/dashboard` route in `web_app.py` renders `templates/dashboard.html`, passing:

- `stats` – from `summary_stats()`.
- `funnel` – from `funnel_metrics()`.
- `paths` – from `session_paths()`.
- `intents` – from `intent_clusters()`.

The HTML shows:

1. KPI cards for basic stats and funnel metrics.
2. Embedded Altair charts in iframes:
  - Top Clicked Products
  - Top Queries
  - Top Query Terms
  - Searches per Hour
  - Term Co-occurrence Heatmap
3. Lists of session paths and intent clusters.

We have seen all the screenshots above that show how we are seeing all those metrics.

## 5. How to Run and Reproduce

1. Install dependencies (Flask, pydantic, pandas, altair, groq, httpagentparser, python-dotenv, etc.).
2. Create a `.env` file (not committed to GitHub) with:

```
Shell
SECRET_KEY=your_flask_secret
SESSION_COOKIE_NAME=irwa_session
DATA_FILE_PATH=../data/fashion_products_dataset_enriched.json
GROQ_API_KEY=your_groq_key
GROQ_MODEL=llama-3.1-8b-instant
DEBUG=True
```

3. Start the web app:

```
Shell
python web_app.py
```

4. Open `http://127.0.0.1:8088` in a browser and issue several queries as described earlier, clicking on some products and waiting a few seconds before starting new searches.
5. Visit `/stats` to see top clicked products and `/dashboard` to see analytics.

## 6. Use of AI Assistance

For Part 4, we used **ChatGPT (OpenAI GPT-5.1 Thinking)** as an assistant to:

- Discuss design options for the RAG improvements and analytics data model.
- Suggest concrete code snippets (especially for prompt engineering, helper scoring, and analytics plots).
- Help draft and structure this written report.

All suggestions were manually reviewed, adapted and tested by us. Any mistakes or omissions in the final code and report remain our responsibility.